

Implementing a PINN with NumPy and Numba (CUDA)

Nikita Sviridov

(translated from the Russian original)

1 Introduction

A PINN (Physics-Informed Neural Network) is a network designed to solve differential equations. The task can be viewed as unsupervised learning, because no labels are needed to train the network. The loss is computed from the differential equation itself, together with the boundary and initial conditions.

The network is trained so that, when the function it represents is substituted into our PDE, BCs and ICs, the deviations from the true values are minimal.

2 The differential equation

I chose the heat equation — a partial differential equation of up to second order:

$$\left\{ \begin{array}{l} \frac{\partial u}{\partial t} = \Delta u + x t^2 \sin y, \quad 0 < x < 1, \ 0 < y < \frac{\pi}{2}, \ t > 0, \\ \frac{\partial u}{\partial x} \Big|_{x=0} = \frac{\partial u}{\partial x} \Big|_{x=1} = 0, \\ u \Big|_{y=0} = 0, \quad \frac{\partial u}{\partial y} \Big|_{y=\pi/2} = 0, \\ u \Big|_{t=0} = 0. \end{array} \right.$$

3 Network architecture

The architecture is a 4-layer multilayer perceptron with a sigmoid activation between layers:

$$\sigma(\sigma(\sigma(XA_1 + B_1)A_2 + B_2)A_3 + B_3)A_4 + B_4,$$

where A_i is the weight matrix of a layer and B_i its bias (added row-wise), and

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

is the sigmoid activation. Weights are initialized with the Xavier method: matrix entries are drawn from a normal distribution $\mathcal{N}(0, 1/\dim_{\text{in}})$.

4 Implementation

There are two stages. The first is writing the training algorithm itself; the second is moving the computation onto CUDA.

4.1 The algorithm

A neural network can be seen as a large parameterized function. Training its parameters, like any model, is based on gradient descent: a loss function is defined, an error is computed from it during training, and then the parameter gradients are obtained. Since the gradient points in the direction of increase, the parameters are shifted along the negative gradient to minimize the error.

Rather than deriving per-layer gradient formulas by hand, I wrote a general form of the network and trained that. By “general form” I mean an automatic-differentiation system — a wrapper class (`tensor`) over an ordinary tensor that records the gradient of every arithmetic operation at each step. This makes it possible to compute the gradient of any node in the computation graph by backpropagation. Because the backward functions are themselves expressed as tensor operations, the gradient graph is also differentiable — which is exactly what allows taking the second-order derivatives (u_{xx} , u_{yy}) required by the equation.

4.2 Parallel computation

To make the network run in reasonable time, the computation is moved onto CUDA. This requires solving the data-allocation problem: training is a dynamic process in which hundreds of constructors and destructors of our class are called at each step. Moving matrices to the GPU during training is inefficient, so a large block of memory is allocated up front and we operate within it.

I wrote a wrapper over the allocated memory, `memory_manager`, to which allocation and reclamation are delegated. When a new tensor is needed, a free region is selected and marked as used; when the destructor is called, the region becomes available again. Region selection uses a greedy fragmentation-minimizing strategy: free intervals are kept sorted by size and the best fit is found in logarithmic time via binary search. On free, the region is merged with adjacent free intervals.

The parallelized functions fall into three categories:

- arithmetic operations;
- memory filling;
- reduce operations.

Arithmetic operations. Implemented: matrix multiplication; elementwise addition, multiplication, division and exponentiation; elementwise $\ln$, $\cos$, $\sin$. The matrix multiply uses a tiled (block) algorithm with shared memory. All binary operators are also defined for the case where the second argument is a scalar.

Memory filling. Implemented: filling a matrix with another matrix or a scalar, taking a slice of a matrix, and writing into a slice.

Reduce operations. Implemented: summation of all elements, summation along columns, and summation along rows.

All computation is implemented in two interchangeable backends: pure NumPy (CPU) and CUDA kernels (Numba). The autodifferentiation graph is identical for both.

5 Training

Loss function. The loss is a sum of three kinds of terms: the first enforces the differential equation, the second the boundary conditions, and the third the initial condition:

$$\text{Loss} = \left\| \frac{\partial u_{\text{pinn}}}{\partial t} - \Delta u_{\text{pinn}} - xt^2 \sin y \right\| + \sum_{i=1}^4 \| \text{BC}_i(u_{\text{pinn}}) \| + \| \text{Init}(u_{\text{pinn}}) \|,$$

where BC_i is the boundary-condition operator and Init the initial-condition operator.

Training parameters. Plain gradient subtraction is not optimal, so a more advanced method — Adam — is used. My implementation is canonical Adam with bias correction of the moments ($\hat{m}$, $\hat{v}$). A learning rate of 0.01 worked best; the number of epochs was set to 2000 as a trade-off between training time and accuracy.

Batch sizes in the loss terms. At each epoch, points are sampled at random from the domain over which the solution is sought. For the PDE error, points are taken from the 3-D space (x, y, t) ; for the ICs and BCs one coordinate is fixed. The PDE error is averaged over 200 points, and each BC/IC error over 50 points.

6 Comparison with the analytic solution

The analytic solution is obtained by separation of variables and expanding the source in a Fourier series in cosines (which satisfy the Neumann conditions in x) and in $\sin y$ (which satisfies the conditions in y). The maximum absolute deviation of the numerical solution from the analytic one does not exceed ≈ 0.01 at $t = 1$.




Figure 1: Left: comparison with the analytic solution (blue — analytic, red — numerical). Right: absolute-error field at $t = 1$.

On the relevance of solving PDEs with PINNs: the approach can compete with other numerical PDE solvers — with library implementations the number of epochs can be increased substantially and the accuracy improved by orders of magnitude.




Figure 2: Training loss over epochs (logarithmic scale).

7 Performance

I compared automatic differentiation running on CUDA versus plain NumPy. At the small matrix sizes used in this architecture (no axis exceeds a few hundred), the GPU is under-utilized, so the dominant cost is not the computation itself but the orchestration: thousands of tiny kernel launches per epoch and an autodifferentiation graph built from Python objects. This also explains why the from-scratch implementation is slower than PyTorch: its dispatch and graph are implemented in C++ and its operations are fused. The bottleneck, therefore, is not the kernel language but the orchestration and the problem size.